

# Dubey

September 2, 2017

```
In [48]: ## In the Below Data Challenge we will Check for the outlier treatment, Missing Value  
        ## after removal of autocorrelation and finally predict the output and the AUC score i
```

```
In [1]: #Importing the Data Set
```

```
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn import preprocessing  
from xgboost.sklearn import XGBClassifier  
from sklearn.model_selection import cross_val_score
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\cross_validation.py:44: DeprecationWarning:  
"This module will be removed in 0.20.", DeprecationWarning)
```

```
In [2]: #Importing the Data
```

```
train = pd.read_csv("D:\\Kaggle Projects\\Blue Owl\\train.csv")  
test = pd.read_csv("D:\\Kaggle Projects\\Blue Owl\\test.csv")
```

```
In [3]: test_raw = test.copy()
```

```
In [4]: test_raw.head(5)
```

```
Out[4]:
```

	age	cost_of_ad	device_type	gender	in_initial_launch_location	income	\
0	34	0.005134	Android	F	1	40376	
1	53	0.005223	desktop	F	1	84511	
2	46	0.004939	laptop	F	0	79322	
3	36	0.004924	Android	F	0	63295	
4	28	0.005146	other	F	1	36170	

	n_drivers	n_vehicles	prior_ins_tenure
0	1	3	7
1	1	1	11
2	1	1	4
3	1	2	0
4	1	3	3

```
In [5]: train.head(5)
```

```
Out [5]:
```

	age	cost_of_ad	device_type	gender	in_initial_launch_location	income	\
0	56	0.005737	iPhone	M		0	62717
1	50	0.004733	desktop	F		0	64328
2	54	0.004129	laptop	M		0	83439
3	16	0.005117	Android	F		0	30110
4	37	0.003635	desktop	M		0	76565


	n_drivers	n_vehicles	prior_ins_tenure	outcome
0	2	1	4	0
1	2	3	2	0
2	1	3	7	0
3	2	3	0	0
4	2	1	5	0

```
In [6]: train.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 10 columns):
age                10000 non-null int64
cost_of_ad         10000 non-null float64
device_type        10000 non-null object
gender             9731 non-null object
in_initial_launch_location 10000 non-null int64
income            10000 non-null int64
n_drivers          10000 non-null int64
n_vehicles         10000 non-null int64
prior_ins_tenure   10000 non-null int64
outcome           10000 non-null int64
dtypes: float64(1), int64(7), object(2)
memory usage: 781.3+ KB
```

```
In [7]: x_age = train['age']
```

```
In [8]: plt.boxplot(x_age)
plt.show()
```




```
In [9]: def transform_diff_grp(dl):  
        if dl > 55 : return 1  
        elif 45 < dl <= 55 : return 2  
        elif 25 < dl <= 45 : return 3  
        elif 15 < dl <= 25 : return 4  
        elif 0 < dl <= 15 : return 5  
  
In [10]: train["age"] = train['age'].map(transform_diff_grp)  
         test["age"] = test['age'].map(transform_diff_grp)  
  
In [11]: age = train['age']  
  
In [12]: plt.hist(age)  
         plt.show()
```



```
In [13]: ad_cost = train['cost_of_ad']
```

```
In [14]: plt.hist(ad_cost)  
plt.show()
```



```
In [15]: train.head(5)
```

```
Out [15]:
```

	age	cost_of_ad	device_type	gender	in_initial_launch_location	income	\
0	1	0.005737	iPhone	M	0	62717	
1	2	0.004733	desktop	F	0	64328	
2	2	0.004129	laptop	M	0	83439	
3	4	0.005117	Android	F	0	30110	
4	3	0.003635	desktop	M	0	76565	

	n_drivers	n_vehicles	prior_ins_tenure	outcome
0	2	1	4	0
1	2	3	2	0
2	1	3	7	0
3	2	3	0	0
4	2	1	5	0

```
In [16]: train['device_type'].unique()
```

```
Out [16]: array(['iPhone', 'desktop', 'laptop', 'Android', 'other'], dtype=object)
```

```
In [17]: train.loc[(train.device_type== 'iPhone') , 'device_type' ] = 0
train.loc[(train.device_type== 'desktop') , 'device_type' ] = 1
train.loc[(train.device_type== 'laptop') , 'device_type' ] = 2
train.loc[(train.device_type== 'Android') , 'device_type' ] = 3
train.loc[(train.device_type== 'other') , 'device_type' ] = 4

test.loc[(test.device_type== 'iPhone') , 'device_type' ] = 0
test.loc[(test.device_type== 'desktop') , 'device_type' ] = 1
test.loc[(test.device_type== 'laptop') , 'device_type' ] = 2
test.loc[(test.device_type== 'Android') , 'device_type' ] = 3
test.loc[(test.device_type== 'other') , 'device_type' ] = 4
```

```
In [18]: test['gender'].unique()
```


```
Out [18]: array(['F', 'M', nan], dtype=object)
```

```
In [19]: test['gender'].fillna(1, inplace=True)
train['gender'].fillna(1, inplace=True)
```

```
In [20]: train.loc[(train.gender== 'M') , 'gender' ] = 1
train.loc[(train.gender== 'F') , 'gender' ] = 0
```


```
In [21]: test.loc[(test.gender== 'M') , 'gender' ] = 1
test.loc[(test.gender== 'F') , 'gender' ] = 0
```

```
In [22]: train.groupby(['gender'])['outcome'].sum().plot(kind="bar")
plt.show()
```



```
In [23]: inc = train['income']
```

```
In [24]: plt.hist(inc)  
plt.show()
```




```
In [25]: train['prior_ins_tenure'].unique()
```

```
Out[25]: array([ 4,  2,  7,  0,  5,  1, 10,  8,  3,  6, 11, 18,  9, 17, 16, 19, 13,
                12, 15, 14, 22, 21, 20], dtype=int64)
```


```
In [26]: ins_ten = train['prior_ins_tenure']
```

```
In [27]: plt.hist(ins_ten)
plt.show()
```



```
In [28]: ins_new = preprocessing.scale(np.sqrt(train['prior_ins_tenure']))
test_ins_new = preprocessing.scale(np.sqrt(test['prior_ins_tenure']))
```


```
In [29]: plt.hist(ins_new)
plt.show()
```



```
In [30]: train['prior_ins_tenure'] = ins_new  
         test['prior_ins_tenure'] = test_ins_new
```

```
In [31]: ins_vec = train['n_drivers']
```

```
In [32]: plt.hist(ins_vec)  
         plt.show()
```





```
In [33]: train['device_type'] = train['device_type'].apply(int)
         train['gender'] = train['gender'].apply(int)
         test['device_type'] = test['device_type'].apply(int)
         test['gender'] = test['gender'].apply(int)
```

```
In [34]: train['n_drivers'].unique()
```

```
Out[34]: array([2, 1], dtype=int64)
```

```
In [35]: # Checking the Correlation Matrix
```

```
Corr_Analysis_Train = train.corr()
```

```
Corr_Analysis_Train.to_csv("D:\\Kaggle Projects\\Blue Owl\\Corr_Analysis_Train.csv",
```

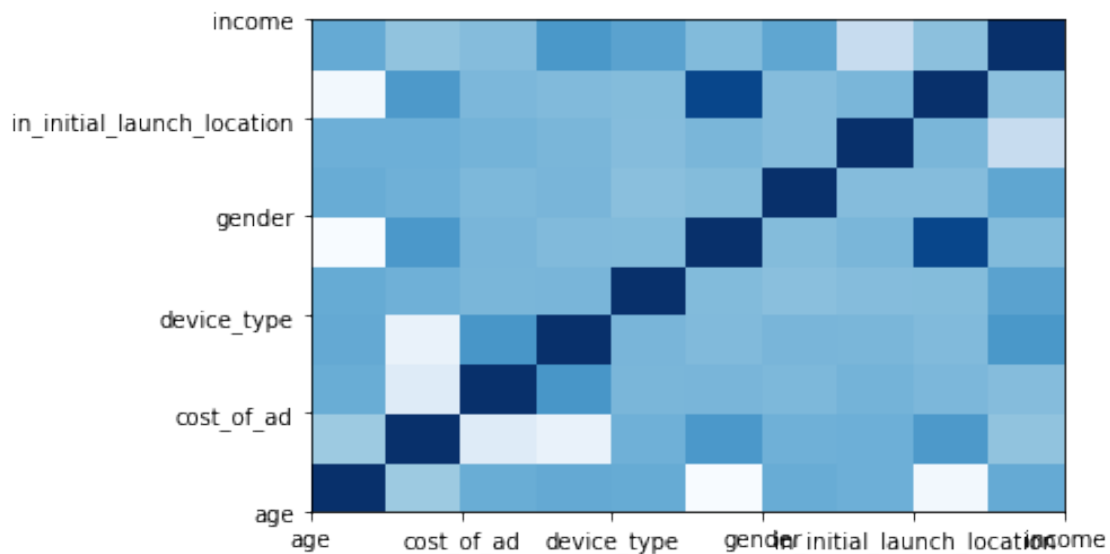
```
fig, ax = plt.subplots()
```

```
ax.pcolor(Corr_Analysis_Train.corr(), cmap=plt.cm.Blues)
```

```
ax.set_xticklabels(Corr_Analysis_Train.columns)
```

```
ax.set_yticklabels(Corr_Analysis_Train.columns)
```

```
plt.show()
```



```
In [36]: y_train = train['outcome']
         train_new = train.drop('outcome', 1)
```

```
In [37]: # Checking the Correlation Matrix
```

```
Corr_Analysis_Train = train_new.corr()
```

```
#Corr_Analysis_Train.to_csv("D:\\Kaggle Projects\\Blue Owl\\Corr_Analysis_Train.csv",
```

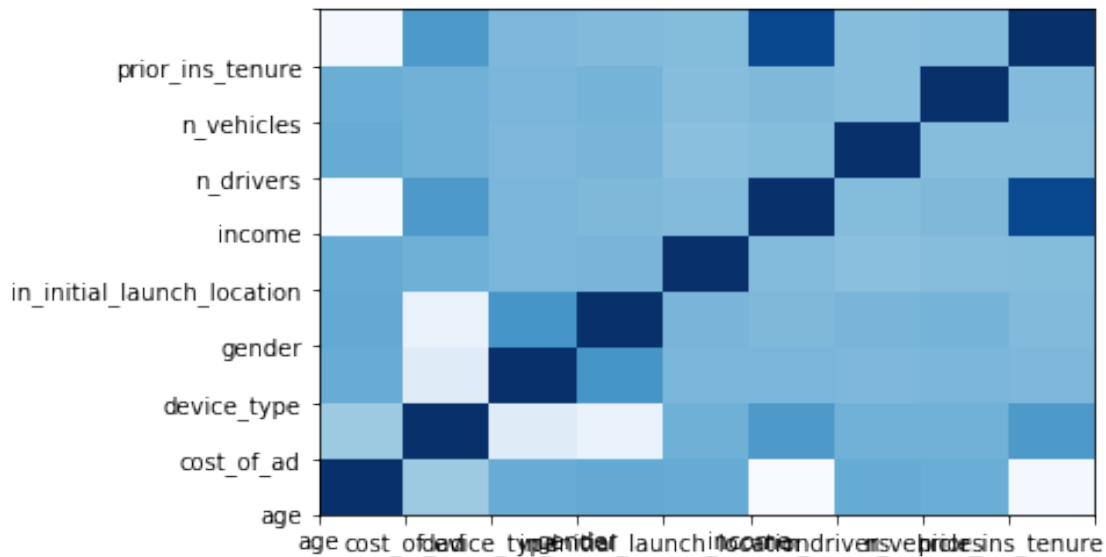
```
fig, ax = plt.subplots()
```

```
ax.pcolor(Corr_Analysis_Train.corr(), cmap=plt.cm.Blues)
```

```

ax.set_xticklabels(Corr_Analysis_Train.columns)
ax.set_yticklabels(Corr_Analysis_Train.columns)
plt.show()

```



```

In [38]: train_new[train_new.dtypes[(train_new.dtypes=="float64")|(train_new.dtypes=="int64")]]
        .index.values].hist(figsize=[11,11])

```

```

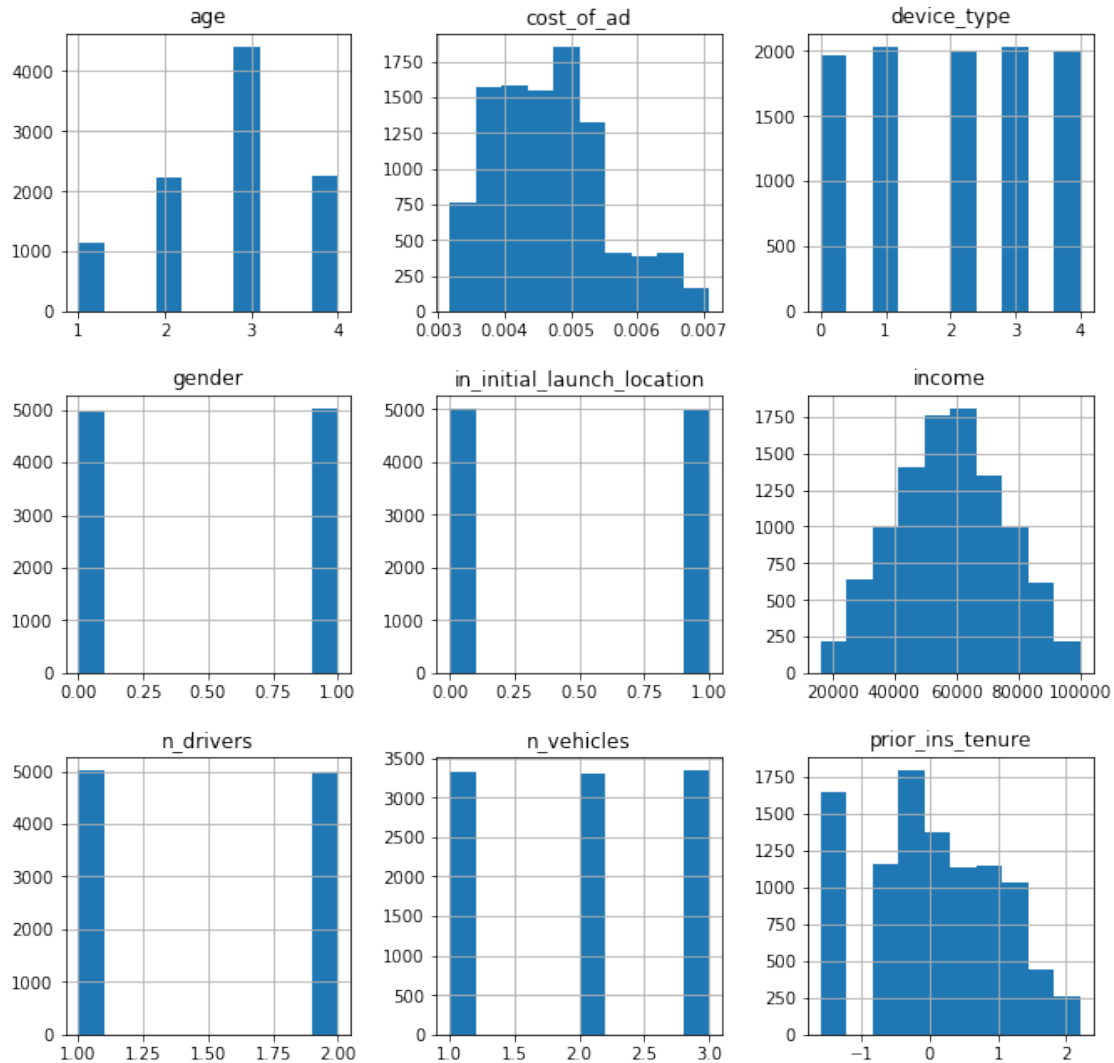
Out[38]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000002066C6AFCF8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x000002066C9EEE48>,
<matplotlib.axes._subplots.AxesSubplot object at 0x000002066C6BDAC8>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x000002066CAB3A20>,
<matplotlib.axes._subplots.AxesSubplot object at 0x000002066CB12BA8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x000002066CB12BE0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x000002066C808E48>,
<matplotlib.axes._subplots.AxesSubplot object at 0x000002066C5E9080>,
<matplotlib.axes._subplots.AxesSubplot object at 0x000002066C7C7BE0>]], dtype=object)

```

```

In [39]: plt.show()

```



```
In [40]: model = XGBClassifier()
         model.fit(train_new, y_train)
```

```
Out[40]: XGBClassifier(base_score=0.5, colsample_bylevel=1, colsample_bytree=1,
                       gamma=0, learning_rate=0.1, max_delta_step=0, max_depth=3,
                       min_child_weight=1, missing=None, n_estimators=100, nthread=-1,
                       objective='binary:logistic', reg_alpha=0, reg_lambda=1,
                       scale_pos_weight=1, seed=0, silent=True, subsample=1)
```

```
In [41]: roc = cross_val_score(model, train_new, y_train, scoring='roc_auc', cv=5)
```

```
In [42]: np.mean(roc)
```

```
Out[42]: 0.84355108689757863
```

```
In [43]: y_pred = model.predict(test)
```

```
In [44]: test_raw.head(5)
```

```
Out [44]:
```

	age	cost_of_ad	device_type	gender	in_initial_launch_location	income	\
0	34	0.005134	Android	F	1	40376	
1	53	0.005223	desktop	F	1	84511	
2	46	0.004939	laptop	F	0	79322	
3	36	0.004924	Android	F	0	63295	
4	28	0.005146	other	F	1	36170	

	n_drivers	n_vehicles	prior_ins_tenure
0	1	3	7
1	1	1	11
2	1	1	4
3	1	2	0
4	1	3	3

```
In [45]: test_raw['outcome'] = y_pred
```

```
In [46]: test_raw.head(15)
```

```
Out [46]:
```

	age	cost_of_ad	device_type	gender	in_initial_launch_location	income	\
0	34	0.005134	Android	F	1	40376	
1	53	0.005223	desktop	F	1	84511	
2	46	0.004939	laptop	F	0	79322	
3	36	0.004924	Android	F	0	63295	
4	28	0.005146	other	F	1	36170	
5	51	0.006242	iPhone	F	0	60520	
6	20	0.003534	desktop	M	0	59324	
7	35	0.004568	Android	F	0	37002	
8	32	0.004713	Android	F	1	45207	
9	33	0.006178	iPhone	F	0	72587	
10	27	0.003350	laptop	M	0	52713	
11	51	0.005774	iPhone	M	0	86333	
12	28	0.003912	desktop	M	1	41397	
13	57	0.003603	desktop	M	0	72786	
14	55	0.003648	Android	M	1	90290	

	n_drivers	n_vehicles	prior_ins_tenure	outcome
0	1	3	7	0
1	1	1	11	0
2	1	1	4	0
3	1	2	0	0
4	1	3	3	0
5	1	1	14	0
6	1	1	0	0
7	2	3	5	0
8	1	2	7	0

9	1	2	6	0
10	2	3	5	0
11	1	1	16	0
12	1	2	4	0
13	2	3	10	0
14	1	1	12	0

```
In [47]: test_raw.to_csv("D:\\Kaggle Projects\\Blue Owl\\FinalResult.csv", index=True , header=
```

```
In [ ]:
```